

CHILDREN'S ALCOHOL CONSUMPTION

Machine Learning

Data Science and Advanced Analytics

Rita Franco
Rodrigo Umbelino
Vitor Manita

m20180080
m20180060
m20180054

01. Abstract

The aim of this project is to build a regressive model that is able to predict, at a macro level, the percentage of children with dangerous alcohol behaviors in a country. In order to achieve this, we applied several techniques of data preprocessing, feature engineering, Machine Learning models and, additionally, the introduction of Genetic Programming within the scope of the project. The focus was on benchmarking different models with different parameters and compare and access the results. The best model should be a model that will have a good generalization ability and small variation of its predictions, allowing us to make rather good conclusions and predictions. Regarding the target variable, we want to predict the alcohol frequency, represented by the proxy feature **alcopops**.

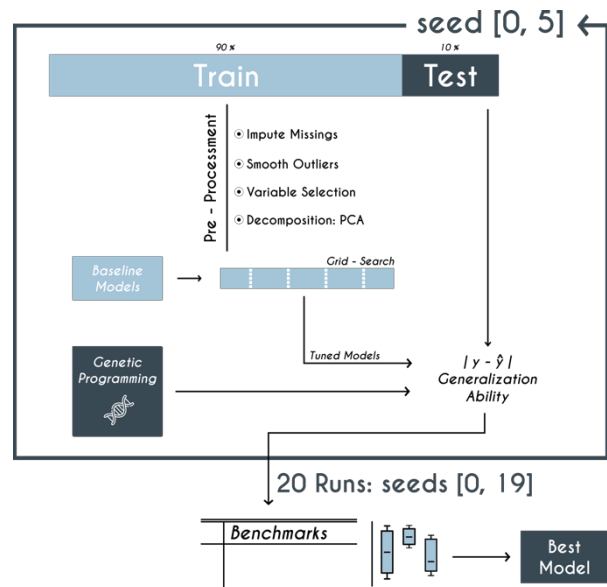


Figure 1. Workflow diagram

02. Explore

Having collected the data from three different sources and three different years, we had to do a lot of preprocessing to get all the data on the same granularity. Our first approach in the exploration phase was to overview the variables and their distribution, briefly discussed on the *Project 2 Description*. After a brief exploration, we could observe which countries presented higher frequencies of alcohol consumption. In *Figure 2*, we present the gender gap in the last collected year (2014). It is also noticeable that boys in Portugal claim they have the 15th highest alcohol frequency consumption between all collected countries.




Figure 2. Children's Alcohol Frequency gender gap (2014)

02.1 Missing Values

Before splitting the data, we need to understand the quality of the data, namely, the missing values in the dataset. As we can see from the Figure 3., we have some 6 features whose missing value percentage is higher than the established 3 % threshold. One of them is our target variables. For the observations with missing values in the target variable we decided to drop them, since they will bring no predictive power to our models. For the remaining features, below and above the threshold, since we have a small dataset, we decided not to drop any observation and fill the missing values.




Figure 3. Missing Values

02.2 Correlations

The next step in the exploration phase is to see how features are correlated amongst themselves. At first sight, in chart 3, we can see that the most noticeable correlation with the alcohol frequency is the year when the data was collected and with the children having his/her first drunk experiences at the ages of 13 years old or younger, with correlation values of -0.43 and 0.51, accordingly. For the remaining independent features, we can also observe a strong correlation between the country GDP per capita and the teenage pregnancy (-0.56), exercise and teenage pregnancy (-0.52) and, the highest dependency, the age of the first drunk experience and if the child has been drunk or not (0.91).




Figure 4. Correlation

03. Modify

For the following chapters and project framework, we had to split the data between Training and Testing. As we have a small dataset, the percentage defined for testing was set at 10%. Furthermore, we decided to perform this action with 5 different seeds, so in the end, we can get the generalization ability of the models in 5 different testing sets and make better conclusions. Since the scope and objective of this project is not to explore all the possibilities in terms of preprocessing of data, we did not experiment a lot of different methods.

For all decisions, we performed all possibilities and benchmarked them, however, in each chapter, we will present the boxplot of the results only depending on that specific decision, over five runs, i.e.: For Imputing Missing values, all other decisions are turned off except for this option.

03.1 Imputing Missing Values

For imputation of Missing Values, we compared the generalization ability of dealing with missing values or just dropping them. It is important to notice that if we drop all observations with missing values in the training set, we would get around half of the dataset size (from around 200 to 100 observations).




Figure 5. Impute Missings Benchmark

03.2 Deal with Outliers

Given that we are handling with data at a macro level, we do not expect to encounter a high number of outliers, since all individual observations will be smoothed by the majority. Having a small of data, we don't want to remove observations classified as outliers, instead, we want to smooth them, hoping that this will improve the predictive ability of the models. As outlier options, we decided to ignore them or to deal with them, by applying a smoothing technique that will, for all features, establish the 5th and 95th percentile and, for values smaller of higher than these thresholds, smooth the data according to the closest of the selected percentiles. *A Priori*, we expect that dealing with outliers, in this case, will not have a huge impact in the results.




Figure 6. Smoothing Outliers Benchmark

04. Feature Selection

Having 21 independent features, another important point to make is whether we will ignore the Curse of Dimensionality or reduce the number of variables, since it can improve the computational effort of the algorithms and, by removing redundant features, avoid multicollinearity issues. However, since the process is a trade-off, we can expect to lose some information. For this method, we performed a Recursive Feature Analysis with a Linear Regressor estimator.




Figure 7. Variable Selection Benchmark

05. Decomposition: PCA

Instead of performing feature selection or, on top of this step, we can also decompose the data into principal components, which are not correlated amongst themselves. For deciding the ideal number of components, since we cannot find a way of looking at the elbow graph automatically in all the runs, we are setting the threshold on the component whose cumulative variance withholds 80% of the explained total variance.




Figure 8. PCA Transformation Benchmark

06. Baseline Models Benchmark

For the baseline models we experimented the following models:

- Random Forest Regressor
- Gradient Boosting Regressor
- Ada Boost Regressor
- Bagging Regressor

We experiment all preprocess option combination, over 5 different runs and seeds, resulting in 56 different outputs. In each run, a grid search with 5 cross validations was performed in order to find the best hyperparameters and use the best estimator of each model. In table 1, the best and worst five models, according to the *Negative Mean Absolute Error*, are shown.

A Priori assumptions

We expect the best generalization ability from the models with missing imputation and outlier smoothing, as well as variable selection.

	Model	Impute Missings	Smooth Outliers	Variable Selection	PCA	NMAE	MAE standard Deviation
Best 5 Models	Gradient Boost	True	True	False	False	-1.286	0.183
	Gradient Boost	True	False	False	False	-1.286	0.182
	Gradient Boost	True	False	True	False	-1.286	0.182
	Bagging	True	False	True	False	-1.320	0.160
	Random Forest	True	True	False	False	-1.330	0.160
Worst 5 Models	Bagging	False	True	True	True	-1.963	0.377
	Random Forest	False	True	True	True	-1.845	0.317
	Gradient Boost	False	True	True	True	-1.842	0.346
	Bagging	True	True	True	True	-1.838	0.374
	Gradient Boost	False	False	True	False	-1.829	0.363

Table 1. Baseline models Benchmark

From Table 1., we can see that, amongst the baseline models (tuned in grid search), Gradient Boost appears to have better results, as well as a rather small variation in its average results. To make a further conclusion about the best decisions to make in preprocessing, we decided to evaluate the generalization ability of the four models grouped by the decisions themselves. For this purpose, we created an id that represents each model according to the choices that were made:

1. **Impute Missings:** *True or False*
2. **Smooth Outliers:** *True or False*
3. **Variable Selection:** *True or False*
4. **PCA:** *True or False*

$\{(T, F); (T, F); (T, F); (T, F)\}$

So, for example, *TTFF* stands for imputing Missings, dealing with outliers but not performing variable selection and decomposition.

From *Figure 9*, we can see that some decisions achieve a higher generalization ability than others. Overall, imputing Missings appears to be a good call. On the other hand, decomposition and variable selection appear to lower the generalization ability of the models, unlike our *a priori* assumptions the most difficult decision to make was whether outliers should be smoothed or not. Since the difference between this choice is not so relevant, we decided to play it safe and perform outlier smoothing, making our decision representation as $\{T, T, F, F\}$.




Figure 9. Pre-process decisions benchmark

06.1 Artificial Data

Having a small amount of data, we wondered if producing artificial data would improve the generalization ability of our models. Having in consideration that we are dealing with a regression problem, the usual packages to generate data such as *SMOTE* did not work, since they are oriented towards a classification problem. Following the same logic, we decided to try our own technique to generate data. We start by generating clusters using the *K-Means* algorithm and, using the centroids, we calculate the distance between the cluster center and each point in that group and then generate a random point between these two coordinates. We tried different number of clusters and proportion of data to generate, presented in *Table 2*. Based on the results, it appears that it is not worth it, in terms of computational effort, generate artificial data, since the payoff from the generalization ability does not change a lot.

The conclusions obtained in this benchmark were applied to the preprocessing of the data imputed in the genetic programming framework.




Figure 10. Normal Data vs Artificial Generated

	Nr. Clusters	Increase Proportion	NMAE
No Artificial	0	0	-1.383
Artificial	10	20%	-1.390
	15	20%	-1.405
	5	20%	-1.413
	15	50%	-1.460
	15	80%	-1.503
	10	50%	-1.510
	5	50%	-1.519
	10	80%	-1.540
	5	80%	-1.603

Table 2. Artificial data Benchmark

07. Genetic Programming

Having done an overall analysis of the baseline models, we explored the Genetic Programming implementation modified package. After some research, we thought and implemented new methods. Their names and description can be found in *Table 3*.

	Name	Description
Selection	Roulette Wheel	Performs the selection of an individual where the probability of selection is proportional to its fitness value
	Stochastic Universal Sampling	Similar to the Roulette Wheel selection but the fitness of the individuals is smoothed by the distances
	Random	Selects one random individual from the population
	Bloat Control	First, it orders the parents by length size and chooses a pool of the n smaller parents. Then, from those, it selects the one with the highest fitness
	Rank	Performs the selection of an individual where the probability of selection is proportional to the rank of its fitness
Crossover	Uniform	Performs crossover between two parents, by swapping with uniform probability same-arity nodes.
	Simple	Similar to the standard crossover, but the donor subtree mandatorily ends in a terminal node, decreasing tree depth in certain occasions.
	2 Tree Crossover	Similar to the crossover implemented in the baseline, however, instead of performing two cuts, it performs 4 cuts and generates 2 subtrees to replace in the original tree
Mutation	Reverse	Changes Function nodes to the opposite function of each node (ex: multiplication is mutated to division).
	Shake	Same concept to point mutation but selects a random subtree and mutates all the nodes within that subtree.
	Graft	Generates a random subtree and selects a random terminal node of an individual. The subtree is then appended to the selected node.
	Swap	Selects two random Function nodes of same arity values from an individual and swaps them.
	Semantical Sig	From the already implemented semantical operator using the sigmoid function, this is the “semantics” version, following the same train of thought as the already implemented hyperbolic tangent operator.

Table 3. Genetic Programming: Implemented methods

08. Genetic Programming Initial Benchmark

For genetic programming we had some big decisions to make, namely, what selectors, crossover and mutations to use. For choosing the best combination, we decided to follow the funnel framework and start by the best selector to use in our final GP model.

Initialization

For a first approach, we tried the different initialization methods available in the baseline, namely *grow*, *full*, *half and half* and *EDDA*. From the results show in *Figure 11.*, all initializations have the same average performance, except for *EDDA*, that looks promising. However, this method presented a real bad result on its first run (seed 0). We decided to proceed with this method and try to prevent this with other operators or fine tuning of the parameters. Furthermore, we also tried the next benchmarking phase with *Half and Half* to compare it do *EDDA*.




Figure 11. Initializations benchmark

Selectors

Having implemented 5 new selection methods, we experimented them with the same parameters used on the previous benchmark. By looking at *Figure 12*, we can see that *EDDA* has a better performance in almost all selectors and that the one outlier situation continues. Despite having similar results to *tournament* selection, *rank* selection appears to have a better generalization ability and less variability so, we continued using this method.

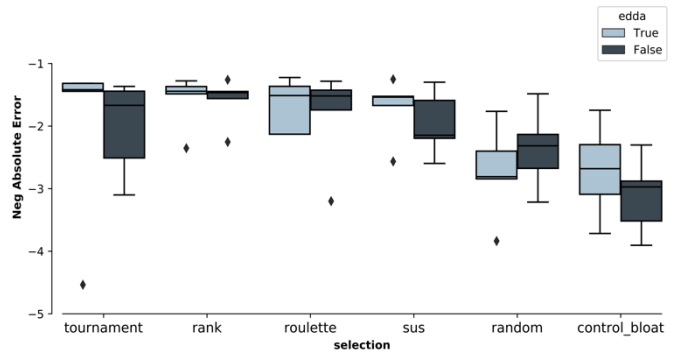


Figure 12. Selection benchmark

Crossover

Having established the initialization method and selector, the following two choices to be made in this initial benchmarking phase was what crossover and mutation method to use. In terms of crossovers, *Uniform*, *Simple* and the baseline crossover were tested, where, presenting similar results, we decided to go with the one provided in the baseline. The 2 Tree Crossover, although it is implemented, it crashed at later runs, so we decided to not benchmark it.




Figure 13. Crossover Benchmarking

Mutation

For mutation operators, we experimented the baseline ones (*Subtre*, *point* and *hoist*), versus the ones that we implemented (*shake*, *graft*, *swap* and *reverse*). For this purpose, we lowered the crossover probability to 0.1 and increases the mutation probability to 0.9. Despite this not making a lot of sense, given the randomness of mutation, our objective was to see what operator presented better initial results.

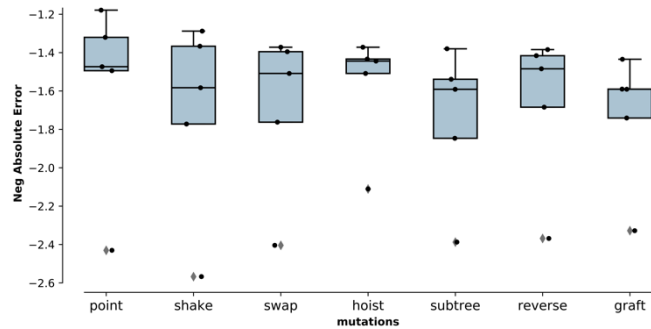


Figure 14. Mutations Benchmark

09. Genetic Programming Fine Tuning

Having the initial and overview benchmark done on genetic programming, we proceeded to the fine tune the algorithm, this is, adjust the hyperparameters such as population size, number of generations, best crossover and mutation probability, stopping criteria and *edda* parameters.

09.1 Population Size

Starting off with the population size, we tried increasing our baseline pool size of 200 to 250, 500, 750 and 1000. The results shown in figure 14 revealed that, in these ranges, for our problem, the generalization ability does not suffer a lot of changes but appears to have better results between a pool size of 250 and 1000, thus, we decided to proceed with a population of 500 individuals.




Figure 15. Population Sizes Benchmark

09.2 Number of Generations

For the number of generations to run, we tried the same intervals experimented on population size, this is: [250, 500, 750, 1000]. Since the results are very similar, we decided to reduce the computational effort and choose 250 generations to run from this point on.




Figure 16. Number of Generations Benchmark

09.3 Crossover and mutations

In terms of crossover and mutation fine tuning, having established the correct crossover and mutation, we want to know to which extent should the probabilities of crossover and mutation affect the generalization ability. Thus, we defined a function that tries crossover with probabilities 25%, 50%, 75% and 100%. For mutation, since the sum of these operators must be equal to




Figure 17. Crossover probability Benchmark

one, it will be the inverse, i.e.: when crossover is at 75%, mutation is at 25%.

09.4 Edda Params

In terms of *EDDA* initialization parameters, we experimented to benchmark the percentage of population calculated with semantical operators, with partitions of 10%, 20%, 30% and 40%, the probability of mutation with 25%, 50%, 75% and 100%. The effect of different levels of maturation was also tested, namely, from 1 to 10 and compared to the baseline provided (5). Lastly, we tried running the genetic programming algorithm with deme sizes of 25, 50, 75 and 100. The results were pretty consistent, but we delimited the following strategy: Running the following epochs with 10% of population being calculated with semantical operators, an *EDDA* mutation probability of 100%, maturation of 9 and a demme size of 50, as settled in the baseline.

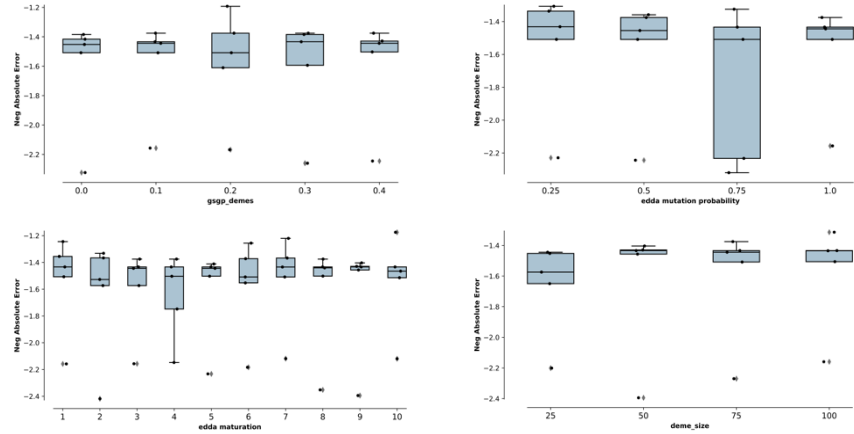


Figure 18. EDDA Parameters Benchmark

10. GP Semantic operators

Before comparing genetic programming to our baseline ensembles, we also tried to run it with semantic operators, with and without semantical computation enabled, that modifies the structure of the trees.

This was a challenging task since we faced some problems in this part, namely the exponential growth of the trees that really diffculted the benchmarking. Due to computational effort, we decided to run this segment with only 10 generations and *Half and Half* initialization. By looking at the results, we can get a clear glance on how semantical computation parameter set to False improves the generalization ability on the crossover provided in the baseline.

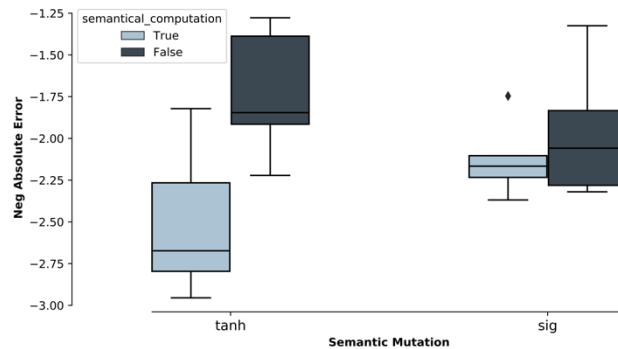


Figure 19. Semantic Operators Benchmark

11. Comparing GP to Ensemble regressors

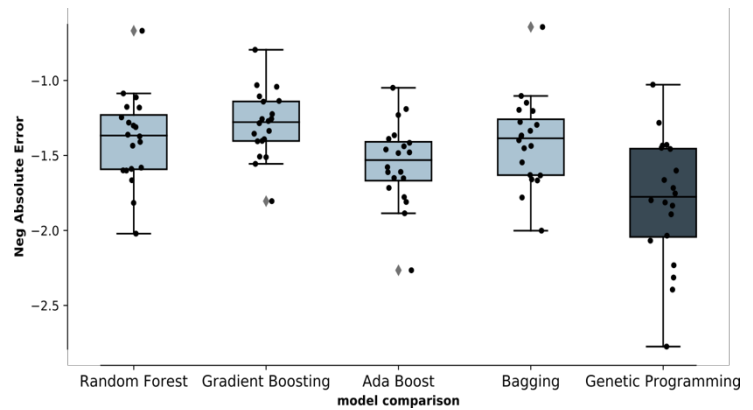
After having all the *Genetic Programming* parameters tuned, we wanted to compare the performance against the baseline ensembles. To have better and more consistent results, we also ran the algorithms more times, totalizing 20 different seeds (runs).

A Priori assumptions

We expect a better performance, in terms of generalization ability and results variation, from a more complex evolutionary algorithm like *Genetic Programming*, when compared to the baseline ensembles.

	Model	Avg NMAE (20 Runs)
Ensembles	Gradient Boosting	-1.290850
	Random Forest	-1.390883
	Bagging	-1.407419
	Ada Boost	-1.553029
GP	Genetic	-1.798638
	Programming	-1.798638

Table 4. 20 Runs Model Average



Contrary to our initial beliefs, what was observed afterwards was that, for this problem, *Genetic Programming* appears to have a lower generalization ability than all baseline ensembles. In the end, as shown in the ensembles benchmark, *Gradient Boosting* regressor reveals to have the best results. Performing a *T-Test* for mean independence, between the results of *Gradient Boosting* and *GP*, we can conclude that, with a *t-score* of 4.79 and a *p-value* of 0.003%, the mean results from these two algorithms, considering the usual confidence thresholds of 1%, 5% and 10%, are indeed different.

12. Conclusions

Overall, this was a very interesting project to tackle, since we got to gather our own data. However, this data collection process was very challenging. We had to read some articles to know what were the most important factors that were influencing children's risky behaviors. Furthermore, besides collecting data from 3 different sources, we had to modify all data from the World Health Organization, transforming each table into just one variable.

After collecting the data, the exploration and preprocessing of data was made. Since the objective of this project was not to explore all methods of pre-processing of data, we decided to follow a simpler train of thought and try a single method for each stage. Contrary to our initial beliefs, variable selection and decomposition did not improve the generalization ability of the models. However, imputing missing values and smoothing outliers appear to, to some extent, help the models achieve a better performance. Then, benchmarking of genetic programming was made, with some new operators and selector techniques implemented. This phase was the most time consuming, since these algorithms are more complex and require more attention. After fine tuning GP, we faced some problems when dealing with semantic GP, that stopped us from better benchmarking these operators.

Lastly, when comparing all algorithms, we reached the conclusion that, for predicting children's alcohol consumption in countries, *Gradient Boosting regressor* appears to be the best choice.

13. Bibliography

- [1] R. Poli and W. Langdon, "On the Search Properties of Different Crossover Operators in Genetic Programming," *Proceedings of Genetic Programming'98*.
- [2] J. Eggermont, "Genetic Programming," *Parallel Problem Solving from Nature - PPSN IX*, 13 9 2006.
- [3] G. F. Guide. [Online]. Available: <https://dces.essex.ac.uk/staff/rpoli/gp-field-guide/22InitialisingthePopulation.html>. [Accessed 5 2019].
- [4] P. D'haeseleer, "Context Preserving Crossover in Genetic Programming".
- [5] A. Swalin, "How to Make Your Machine Learning Models Robust to Outliers," 31 5 2018. [Online]. Available: <https://heartbeat.fritz.ai/how-to-make-your-machine-learning-models-robust-to-outliers-44d404067d07>. [Accessed 5 2019].